


How a question gets an answer

What happens when someone asks Agentic RAG a question. The **API server** works with the **AI models** directly, while you wait.



Nothing relevant found?

If the review (step 5) finds no passage that answers the question, steps 8–9 are skipped. You get **"I couldn't find this in the documents"**, because the AI is never asked to guess.

Agent mode: the AI plans its own searches

Instead of searching once with your words, the chat AI decides **what to search for** and can search **several times**, for example once per part of a two-part question, before writing the answer. Each search is shown live.

The same journey, in plain language

A question takes about 7–12 seconds on our machine, and the answer starts appearing after about 5 seconds.

STEP	WHAT HAPPENS	WHY IT MATTERS	TIME
1–2	You ask; the API server checks your access key.	Only people with a key can ask.	instant
3	AI turns the question into a meaning fingerprint.	So it can be compared with the fingerprints of every passage.	~0.3 s
4	The database searches by meaning and by exact words.	Meaning finds paraphrases; exact words find names, codes and numbers.	< 0.1 s
5	AI reviews the 10 best candidates.	Search is fast but rough; the review checks that a passage really answers the question.	~3–5 s
6–7	The best passages are kept, numbered and shown to you.	You see what the answer will be based on, before it is written.	instant
8	AI writes the answer from those passages only.	Every claim points to its source, like [1].	~3–5 s, streamed
9–10	Citations are checked; you can click each one.	Anyone can verify where a statement came from.	instant

Asking vs. uploading

Uploading is handled in the background by a **worker**, and nobody waits for it. **Asking** is handled by the **API server** itself, because you are waiting for the answer. The same fingerprint model is used in both, so questions and passages can be compared.

Three AI moments, two models

1. **The translator** (step 3, qwen3-embedding) turns your question into a meaning fingerprint.
2. **The reviewer** (step 5, gpt-oss) reads each candidate passage and grades whether it answers the question.
3. **The writer** (step 8, gpt-oss) writes the answer, citing the passages it used.

Guards against made-up answers

- The writer may use **only** the numbered passages, not its general knowledge.
- No relevant passages means no AI answer, just an honest “couldn’t find it”.
- Every citation is checked against the passages that were provided.
- Instructions hidden inside documents are ignored: passages are treated as data, not commands.

Ask or Agent?

Ask: one search with your words, so it is quick and predictable. Best for direct questions.

Agent: the AI plans its own searches. Best for questions with several parts.